

Machine Learning Project

Prabal Ghosh

MSc Data Science & Artificial Intelligence

Université Côte d'Azur

Machine Learning

Date- 16/12/2023

Context: Climate Change and Building Energy Efficiency

Climate change stands as an urgent, multifaceted global issue significantly impacting energy policy and infrastructure. Addressing climate change involves both mitigation (reducing greenhouse gas emissions) and adaptation (preparing for inevitable consequences). Mitigation efforts require changes across electricity systems, transportation, buildings, industries, and land use.

According to a report by the International Energy Agency (IEA), the life cycle of buildings, from construction to demolition, accounted for 37% of global CO₂ emissions associated with energy and processes in 2020. However, there exists substantial potential to decrease buildings' energy consumption by integrating easily implementable solutions with cutting-edge strategies.

For instance, renovated buildings have shown the capability to reduce heating and cooling energy needs by 50-90%. Moreover, many of these energy efficiency measures yield overall cost savings and additional benefits, such as providing occupants with cleaner air. Achieving this potential is viable while upholding the services offered by the buildings.

The Dataset and Challenge

Accurate predictions of energy consumption for a building based on its characteristics are crucial for policymakers to effectively target renovation efforts and maximize emissions reductions.

The dataset utilized originates from the Lawrence Berkeley National Laboratory (Berkeley Lab). The task at hand involves analyzing variations in building energy efficiency to construct one or more predictive models for estimating the energy consumption of buildings.

The provided data encompasses descriptions of building characteristics along with climatic and meteorological variables specific to the regions where these buildings are situated.

Dataset Description

The dataset comprises around 100k observations gathered over a span of 7 years across various locations, focusing on building energy usage.

It encompasses:

- Building characteristics (e.g., floor area, installation type, etc.).
- Meteorological data specific to each building's location (e.g., mean annual temperature, total annual precipitation, etc.).
- Energy consumption details for each building within a given year.

Each row in the dataset corresponds to a singular building observed in a particular year.

Your objective is to predict the Site Energy Use Intensity (EUI) for each row, leveraging the building characteristics and the meteorological data associated with the building's location.

Evaluation Metrics: Negative Root Mean Square Error

Features

- **id:** Building ID
- **Year_Factor:** Anonymized year when weather and energy usage factors were observed
- **State_Factor:** Anonymized state in which the building is located
- **building_class:** Building classification
- **facility_type:** Building usage type
- **floor_area:** Floor area (in square feet) of the building
- **year_built:** Year in which the building was constructed
- **energy_star_rating:** The Energy Star rating of the building
- **ELEVATION:** Elevation of the building location
- **january_min_temp:** Minimum temperature in January (in Fahrenheit) at the location of the building
- **january_avg_temp:** Average temperature in January (in Fahrenheit) at the location of the building
- **january_max_temp:** Maximum temperature in January (in Fahrenheit) at the location of the building
- **cooling_degree_days:** Cooling degree day for a given day, representing the number of degrees where the daily average temperature exceeds 65 degrees Fahrenheit. The monthly sum produces an annual total at the building's location.
- **heating_degree_days:** Heating degree day for a given day, representing the number of degrees where the daily average temperature falls under 65 degrees Fahrenheit. The monthly sum produces an annual total at the building's location.
- **precipitation_inches:** Annual precipitation in inches at the location of the building
- **snowfall_inches:** Annual snowfall in inches at the location of the building
- **snowdepth_inches:** Annual snow depth in inches at the location of the building
- **avg_temp:** Average temperature over a year at the location of the building
- **days_below_30F:** Total number of days below 30 degrees Fahrenheit at the location of the building
- **days_below_20F:** Total number of days below 20 degrees Fahrenheit at the location of the building

- **days_below_10F:** Total number of days below 10 degrees Fahrenheit at the location of the building
- **days_below_0F:** Total number of days below 0 degrees Fahrenheit at the location of the building
- **days_above_80F:** Total number of days above 80 degrees Fahrenheit at the location of the building
- **days_above_90F:** Total number of days above 90 degrees Fahrenheit at the location of the building
- **days_above_100F:** Total number of days above 100 degrees Fahrenheit at the location of the building
- **days_above_110F:** Total number of days above 110 degrees Fahrenheit at the location of the building
- **direction_max_wind_speed:** Wind direction for maximum wind speed at the location of the building, given in 360-degree compass point directions (e.g., 360 = north, 180 = south, etc.).
- **direction_peak_wind_speed:** Wind direction for peak wind gust speed at the location of the building, given in 360-degree compass point directions (e.g., 360 = north, 180 = south, etc.).
- **max_wind_speed:** Maximum wind speed at the location of the building
- **days_with_fog:** Number of days with fog at the location of the building

Target

- **site_eui:** Site Energy Usage Intensity is the amount of heat and electricity consumed by a building as reflected in utility bills.

Your Job

Your task involves demonstrating your knowledge gained from the machine learning course by constructing an sklearn pipeline to perform the following:

1. Data Pre-processing:

- Correctly pre-process the data based on its category. The specific steps for data pre-processing are not provided here as they are covered in various courses.

2. Hyper-parameter Tuning:

- Choose appropriate hyper-parameters for the selected models. While the specific hyper-parameters are not mentioned here (as they have been covered in the course), your task involves identifying and tuning the right hyper-parameters for the models.

An initial Exploratory Data Analysis (EDA) is crucial to understand the nature of each feature within the dataset.

CAUTION:

- Utilize only the concepts learned in the course. If you include any additional concepts, ensure you understand and can explain them thoroughly.

- Avoid blindly copying code from the internet or generated by ChatGPT without comprehension. It's better to have a concise and fully understood piece of work rather than an extensive application of techniques without understanding their functionality or usefulness.

You'll need to present a written notebook (with comments) that justifies your choices for different stages of the pipeline.

Solution

I referred this book for data preprocessing- [Click here to download data preprocessing book](#)

Libraries

```
In [1]: import numpy as np
import pandas as pd

import matplotlib.pyplot as plt
import seaborn as sns
plt.style.use('bmh')
sns.set_style({'axes.grid':False})
%matplotlib inline
```

```
In [2]: import sklearn
from sklearn.decomposition import PCA

from sklearn.pipeline import Pipeline
from sklearn.compose import make_column_selector as selector
from sklearn.preprocessing import OneHotEncoder, StandardScaler, MinMaxScaler
from sklearn.model_selection import train_test_split
from sklearn.impute import SimpleImputer
from sklearn.base import BaseEstimator, TransformerMixin
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import make_pipeline
from sklearn import set_config
from sklearn.metrics import mean_squared_error

from imblearn.pipeline import Pipeline as ImbPipeline

from sklearn.model_selection import GridSearchCV, RandomizedSearchCV
```

```
In [3]: from sklearn.svm import SVR
from sklearn import linear_model

from sklearn.linear_model import LinearRegression, Ridge, RidgeCV, Lasso, LassoCV, ElasticNet
from sklearn.tree import DecisionTreeRegressor, plot_tree

from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import r2_score
```


Dataset

```
In [4]: energy_df = pd.read_csv("C:\\Users\\praba\\Desktop\\uca1\\M1\\ML\\final project\\data\\energy_df.csv")
energy_df.tail(5)
```

```
Out[4]:
```

	Year_Factor	State_Factor	building_class	facility_type	floor_area	year_built	energy_st
75752	6	State_11	Commercial	Office_Uncategorized	20410.0	1995.0	
75753	6	State_11	Residential	5plus_Unit_Building	40489.0	1910.0	
75754	6	State_11	Commercial	Commercial_Other	28072.0	1917.0	
75755	6	State_11	Commercial	Commercial_Other	53575.0	2012.0	
75756	6	State_11	Residential	2to4_Unit_Building	23888.0	1974.0	

5 rows × 64 columns



```
In [5]: energy_df.shape
```

```
Out[5]: (75757, 64)
```

```
In [6]: '''All id values are unique which are irrelevant for model training'''
energy_df['id'].unique().shape
```

```
Out[6]: (75757,)
```

```
In [7]: energy_df['Year_Factor'].unique()
```

```
Out[7]: array([1, 2, 3, 4, 5, 6], dtype=int64)
```

There are a total of six different values for the 'Year_Factor' feature in this dataset. Therefore, I will use one-hot encoding for these values, considering them as distinct levels or categorical entries.

Checking NAN data

I am checking for missing data because if I have the model can't learn properly, since some of the data would have values that don't represent the reality. Generally we use inputers to replace the missing data with statistical measures such as the mean/median for numerical columns and the mode for categorical columns, but first we need to check if there is missing data at all.

```
In [8]: energy_df.describe()
```

```
Out[8]:
```

	Year_Factor	floor_area	year_built	energy_star_rating	ELEVATION	january_min_temp
count	75757.000000	7.575700e+04	73920.000000	49048.000000	75757.000000	75757.000000
mean	4.367755	1.659839e+05	1952.306764	61.048605	39.506323	11.432343
std	1.471441	2.468758e+05	37.053619	28.663683	60.656596	9.381027
min	1.000000	9.430000e+02	0.000000	0.000000	-6.400000	-19.000000
25%	3.000000	6.237900e+04	1927.000000	40.000000	11.900000	6.000000
50%	5.000000	9.136700e+04	1951.000000	67.000000	25.000000	11.000000

	Year_Factor	floor_area	year_built	energy_star_rating	ELEVATION	january_min_temp
75%	6.000000	1.660000e+05	1977.000000	85.000000	42.700000	13.000000
max	6.000000	6.385382e+06	2015.000000	100.000000	1924.500000	49.000000

8 rows × 61 columns

```
In [9]: '''Nan value is present in the dataset'''
energy_df.isnull().any().any()
```

Out[9]: True

```
In [10]: # a = energy_df_id.isna().sum()
# for i in a:
#     if i !=0:
#         print(i)
```

```
In [11]: energy_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 75757 entries, 0 to 75756
Data columns (total 64 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Year_Factor                          75757 non-null  int64
1   State_Factor                        75757 non-null  object
2   building_class                      75757 non-null  object
3   facility_type                      75757 non-null  object
4   floor_area                          75757 non-null  float64
5   year_built                          73920 non-null  float64
6   energy_star_rating                  49048 non-null  float64
7   ELEVATION                          75757 non-null  float64
8   january_min_temp                    75757 non-null  int64
9   january_avg_temp                    75757 non-null  float64
10  january_max_temp                    75757 non-null  int64
11  february_min_temp                   75757 non-null  int64
12  february_avg_temp                   75757 non-null  float64
13  february_max_temp                   75757 non-null  int64
14  march_min_temp                      75757 non-null  int64
15  march_avg_temp                      75757 non-null  float64
16  march_max_temp                      75757 non-null  int64
17  april_min_temp                     75757 non-null  int64
18  april_avg_temp                      75757 non-null  float64
19  april_max_temp                      75757 non-null  int64
20  may_min_temp                        75757 non-null  int64
21  may_avg_temp                        75757 non-null  float64
22  may_max_temp                        75757 non-null  int64
23  june_min_temp                       75757 non-null  int64
24  june_avg_temp                       75757 non-null  float64
25  june_max_temp                       75757 non-null  int64
26  july_min_temp                       75757 non-null  int64
27  july_avg_temp                       75757 non-null  float64
28  july_max_temp                       75757 non-null  int64
29  august_min_temp                     75757 non-null  int64
30  august_avg_temp                     75757 non-null  float64
31  august_max_temp                     75757 non-null  int64
32  september_min_temp                  75757 non-null  int64
```

```

33 september_avg_temp      75757 non-null float64
34 september_max_temp      75757 non-null int64
35 october_min_temp        75757 non-null int64
36 october_avg_temp        75757 non-null float64
37 october_max_temp        75757 non-null int64
38 november_min_temp       75757 non-null int64
39 november_avg_temp       75757 non-null float64
40 november_max_temp       75757 non-null int64
41 december_min_temp       75757 non-null int64
42 december_avg_temp       75757 non-null float64
43 december_max_temp       75757 non-null int64
44 cooling_degree_days      75757 non-null int64
45 heating_degree_days     75757 non-null int64
46 precipitation_inches    75757 non-null float64
47 snowfall_inches         75757 non-null float64
48 snowdepth_inches        75757 non-null int64
49 avg_temp                75757 non-null float64
50 days_below_30F          75757 non-null int64
51 days_below_20F          75757 non-null int64
52 days_below_10F          75757 non-null int64
53 days_below_0F           75757 non-null int64
54 days_above_80F          75757 non-null int64
55 days_above_90F          75757 non-null int64
56 days_above_100F         75757 non-null int64
57 days_above_110F         75757 non-null int64
58 direction_max_wind_speed 34675 non-null float64
59 direction_peak_wind_speed 33946 non-null float64
60 max_wind_speed          34675 non-null float64
61 days_with_fog           29961 non-null float64
62 site_eui                75757 non-null float64
63 id                     75757 non-null int64
dtypes: float64(24), int64(37), object(3)
memory usage: 37.0+ MB

```

```
In [12]: # energy_df.id.isna().sum()
```


```
In [13]: '''Get the features with their corresponding total number of missing values.'''
missing_columns = len(energy_df) - energy_df.loc[:, np.sum(energy_df.isnull())>0].columns
missing_columns
```

```
Out[13]: year_built      1837
energy_star_rating    26709
direction_max_wind_speed 41082
direction_peak_wind_speed 41811
max_wind_speed        41082
days_with_fog         45796
dtype: int64
```

heatmap plot to see the null values in dataset

```
In [14]: '''heatmap plot to see the null values in dataset'''

plt.figure(figsize=(25,25))
sns.heatmap(energy_df.isnull())
plt.show()
```



Columns with Null Values

The following six columns contain null values:

- year_built
- energy_star_rating
- direction_max_wind_speed
- direction_peak_wind_speed
- max_wind_speed
- days_with_fog

Dropping Columns with Null Values exceeding 40,000 instances

Upon inspecting the dataset, it was found that the following columns:

'direction_max_wind_speed', 'direction_peak_wind_speed', 'max_wind_speed', and

'days_with_fog' contain a substantial number of null values, each exceeding 40,000 instances, which represents more than 50% of the data in these columns.

Therefore, due to the significant presence of missing values, I have decided to drop these columns from the dataset. This step ensures a more robust and accurate analysis by excluding columns with inadequate or unreliable data.

```
In [15]: '''All id values are unique which are irrelevant for model training. That's why I am
energy_df = energy_df.loc[:, energy_df.columns != "id"]
# energy_df.head(3)

In [16]: """'direction_max_wind_speed', 'direction_peak_wind_speed', 'max_wind_speed', and 'd
energy_df.drop(['direction_max_wind_speed', 'direction_peak_wind_speed', 'max_wind_spe

In [17]: energy_df.shape

Out[17]: (75757, 59)
```

Null values in 'year_built' column was just approx 1% of the whole data so i will fill those using mean

```
In [18]: energy_df["year_built"].isna().sum()

Out[18]: 1837

In [19]: # energy_df['year_built'].mode()

In [20]: # np.mean(energy_df['year_built'])

In [21]: # energy_df.dropna(subset=['year_built'], inplace=True, axis=0)
# energy_df.shape

In [22]: # Fill NaN values in a specific column with the mean
mean_value_year = energy_df['year_built'].mean()
energy_df['year_built'].fillna(mean_value_year, inplace=True)

In [23]: '''Get the features (with missing values) and their corresponding missing values.'''
missing_columns = len(energy_df) - energy_df.loc[:, np.sum(energy_df.isnull())>0].co
missing_columns

Out[23]: energy_star_rating    26709
dtype: int64


In [24]: # Fill NaN values in a specific column with the mean
mean_value = energy_df['energy_star_rating'].mean()
```

In column "energy_star_rating" 25472 null values present. So I am filling those using the mean value of the column

```
energy_df['energy_star_rating'].fillna(mean_value, inplace=True)
```

```
In [25]: '''heatmap plot to see the null values in dataset=====>>> Here no null value present

plt.figure(figsize=(10,10))
sns.heatmap(energy_df.isnull())
plt.show()
```



```
In [26]: """Now there is no nan value present in the dataset."""
energy_df.isnull().sum().sum()
```

```
Out[26]: 0
```

```
In [27]: # features = energy_df.columns.tolist()
# print("The columns in dataset:----- ",features)
# type(features)
```

There are a total of six different values for the 'Year_Factor' feature in this dataset. Therefore, I will use one-hot encoding for these values, considering them as distinct levels as categorical entries.

```
In [28]: '''Numerical and categorical columns are separated and Year_Factor is chosen as cat
import pandas as pd

features = energy_df.columns.tolist()

Num_features = [feature for feature in features if
                 energy_df[feature].dtype != 'object' and
                 feature != 'Year_Factor']
Cat_features = [feature for feature in features if
                energy_df[feature].dtype == object or feature == 'Year_Factor' ]

print("Num_features: \n",Num_features)
print("\n")

print("Cat_features: \n",Cat_features)
```

```
Num_features:
['floor_area', 'year_built', 'energy_star_rating', 'ELEVATION', 'january_min_temp',
'january_avg_temp', 'january_max_temp', 'february_min_temp', 'february_avg_temp', 'february_max_temp', 'march_min_temp', 'march_avg_temp', 'march_max_temp', 'april_min_temp', 'april_avg_temp', 'april_max_temp', 'may_min_temp', 'may_avg_temp', 'may_max_temp', 'june_min_temp', 'june_avg_temp', 'june_max_temp', 'july_min_temp', 'july_avg_temp', 'july_max_temp', 'august_min_temp', 'august_avg_temp', 'august_max_temp', 'september_min_temp', 'september_avg_temp', 'september_max_temp', 'october_min_temp', 'october_avg_temp', 'october_max_temp', 'november_min_temp', 'november_avg_temp', 'november_max_temp', 'december_min_temp', 'december_avg_temp', 'december_max_temp', 'cooling_degree_days', 'heating_degree_days', 'precipitation_inches', 'snowfall_inches', 'snowdepth_inches', 'avg_temp', 'days_below_30F', 'days_below_20F', 'days_below_10F', 'days_below_0F', 'days_above_80F', 'days_above_90F', 'days_above_100F', 'days_above_110F', 'site_eui']
```

```
Cat_features:
['Year_Factor', 'State_Factor', 'building_class', 'facility_type']
```


```
In [29]: energy_df.hist(figsize=[30,30])
```

```
Out[29]: array([[<AxesSubplot:title={'center':'Year_Factor'}>,
<AxesSubplot:title={'center':'floor_area'}>,
<AxesSubplot:title={'center':'year_built'}>,
<AxesSubplot:title={'center':'energy_star_rating'}>,
<AxesSubplot:title={'center':'ELEVATION'}>,
<AxesSubplot:title={'center':'january_min_temp'}>,
<AxesSubplot:title={'center':'january_avg_temp'}>],
[<AxesSubplot:title={'center':'january_max_temp'}>,
<AxesSubplot:title={'center':'february_min_temp'}>,
<AxesSubplot:title={'center':'february_avg_temp'}>,
<AxesSubplot:title={'center':'february_max_temp'}>,
<AxesSubplot:title={'center':'march_min_temp'}>,
<AxesSubplot:title={'center':'march_avg_temp'}>,
<AxesSubplot:title={'center':'march_max_temp'}>],
[<AxesSubplot:title={'center':'april_min_temp'}>,
<AxesSubplot:title={'center':'april_avg_temp'}>,
<AxesSubplot:title={'center':'april_max_temp'}>,
<AxesSubplot:title={'center':'may_min_temp'}>,
<AxesSubplot:title={'center':'may_avg_temp'}>],
[<AxesSubplot:title={'center':'may_max_temp'}>,
<AxesSubplot:title={'center':'june_min_temp'}>,
<AxesSubplot:title={'center':'june_avg_temp'}>,
<AxesSubplot:title={'center':'june_max_temp'}>,
<AxesSubplot:title={'center':'july_min_temp'}>,
<AxesSubplot:title={'center':'july_avg_temp'}>,
<AxesSubplot:title={'center':'july_max_temp'}>,
<AxesSubplot:title={'center':'august_min_temp'}>,
<AxesSubplot:title={'center':'august_avg_temp'}>,
<AxesSubplot:title={'center':'august_max_temp'}>,
<AxesSubplot:title={'center':'september_min_temp'}>,
<AxesSubplot:title={'center':'september_avg_temp'}>,
<AxesSubplot:title={'center':'september_max_temp'}>,
<AxesSubplot:title={'center':'october_min_temp'}>,
<AxesSubplot:title={'center':'october_avg_temp'}>,
<AxesSubplot:title={'center':'october_max_temp'}>,
<AxesSubplot:title={'center':'november_min_temp'}>,
<AxesSubplot:title={'center':'november_avg_temp'}>,
<AxesSubplot:title={'center':'november_max_temp'}>,
<AxesSubplot:title={'center':'december_min_temp'}>,
<AxesSubplot:title={'center':'december_avg_temp'}>,
<AxesSubplot:title={'center':'december_max_temp'}>,
<AxesSubplot:title={'center':'cooling_degree_days'}>,
<AxesSubplot:title={'center':'heating_degree_days'}>,
<AxesSubplot:title={'center':'precipitation_inches'}>,
<AxesSubplot:title={'center':'snowfall_inches'}>,
<AxesSubplot:title={'center':'snowdepth_inches'}>,
<AxesSubplot:title={'center':'avg_temp'}>,
<AxesSubplot:title={'center':'days_below_30F'}>,
<AxesSubplot:title={'center':'days_below_20F'}>,
<AxesSubplot:title={'center':'days_below_10F'}>,
<AxesSubplot:title={'center':'days_below_0F'}>,
<AxesSubplot:title={'center':'days_above_80F'}>,
<AxesSubplot:title={'center':'days_above_90F'}>,
<AxesSubplot:title={'center':'days_above_100F'}>,
<AxesSubplot:title={'center':'days_above_110F'}>,
<AxesSubplot:title={'center':'site_eui'}>])
```

```

<AxesSubplot:title={'center':'may_max_temp'}>,
<AxesSubplot:title={'center':'june_min_temp'}>],
[<AxesSubplot:title={'center':'june_avg_temp'}>,
<AxesSubplot:title={'center':'june_max_temp'}>,
<AxesSubplot:title={'center':'july_min_temp'}>,
<AxesSubplot:title={'center':'july_avg_temp'}>,
<AxesSubplot:title={'center':'july_max_temp'}>,
<AxesSubplot:title={'center':'august_min_temp'}>,
<AxesSubplot:title={'center':'august_avg_temp'}>],
[<AxesSubplot:title={'center':'august_max_temp'}>,
<AxesSubplot:title={'center':'september_min_temp'}>,
<AxesSubplot:title={'center':'september_avg_temp'}>,
<AxesSubplot:title={'center':'september_max_temp'}>,
<AxesSubplot:title={'center':'october_min_temp'}>,
<AxesSubplot:title={'center':'october_avg_temp'}>,
<AxesSubplot:title={'center':'october_max_temp'}>],
[<AxesSubplot:title={'center':'november_min_temp'}>,
<AxesSubplot:title={'center':'november_avg_temp'}>,
<AxesSubplot:title={'center':'november_max_temp'}>,
<AxesSubplot:title={'center':'december_min_temp'}>,
<AxesSubplot:title={'center':'december_avg_temp'}>,
<AxesSubplot:title={'center':'december_max_temp'}>,
<AxesSubplot:title={'center':'cooling_degree_days'}>],
[<AxesSubplot:title={'center':'heating_degree_days'}>,
<AxesSubplot:title={'center':'precipitation_inches'}>,
<AxesSubplot:title={'center':'snowfall_inches'}>,
<AxesSubplot:title={'center':'snowdepth_inches'}>,
<AxesSubplot:title={'center':'avg_temp'}>,
<AxesSubplot:title={'center':'days_below_30F'}>,
<AxesSubplot:title={'center':'days_below_20F'}>],
[<AxesSubplot:title={'center':'days_below_10F'}>,
<AxesSubplot:title={'center':'days_below_0F'}>,
<AxesSubplot:title={'center':'days_above_80F'}>,
<AxesSubplot:title={'center':'days_above_90F'}>,
<AxesSubplot:title={'center':'days_above_100F'}>,
<AxesSubplot:title={'center':'days_above_110F'}>,
<AxesSubplot:title={'center':'site_eui'}>]], dtype=object)

```



Boxplot to check outliers

In [30]:

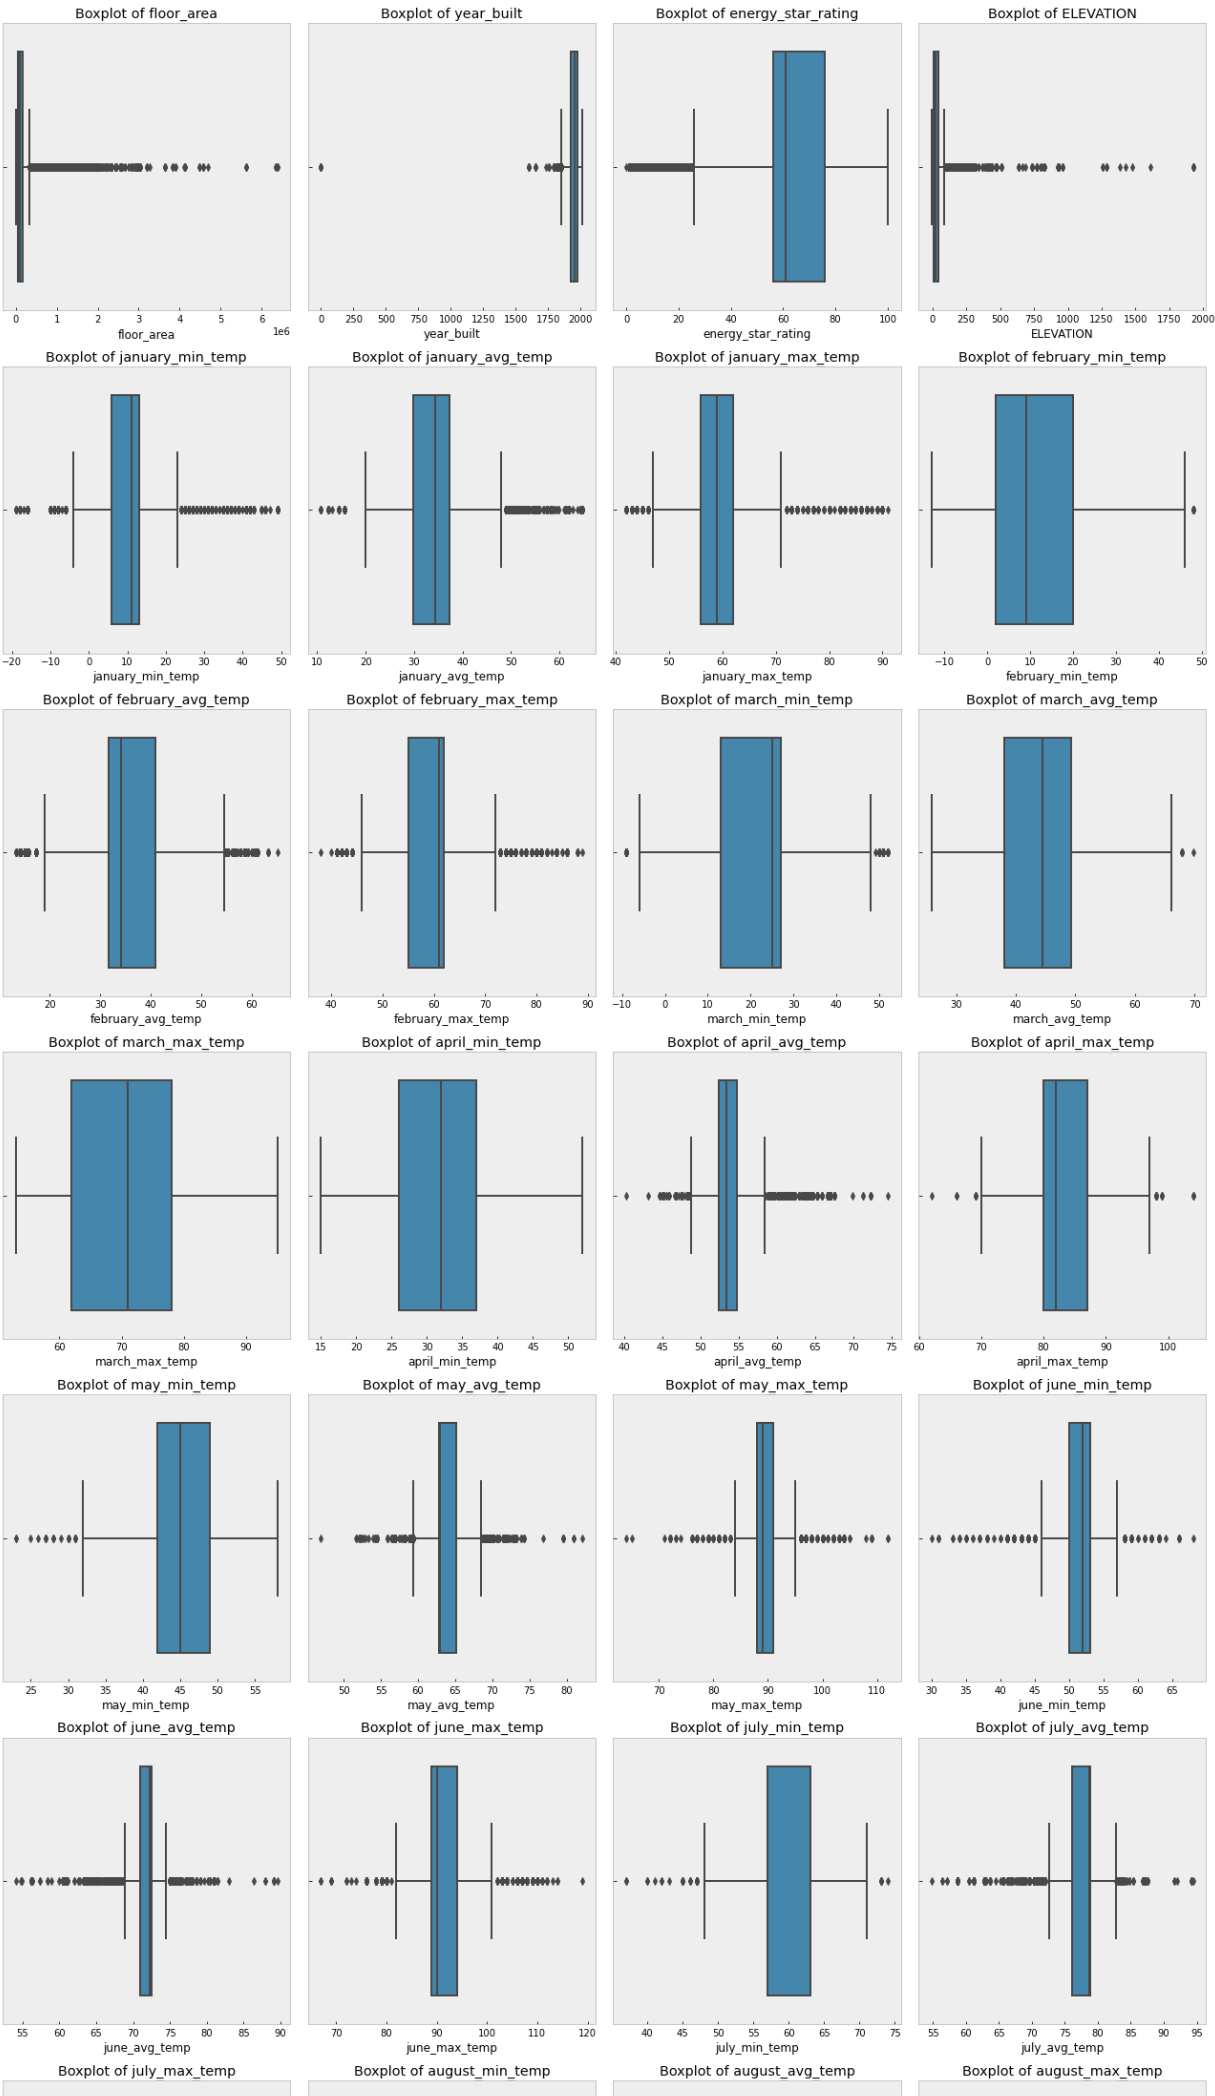
```
import matplotlib.pyplot as plt
import seaborn as sns
import math

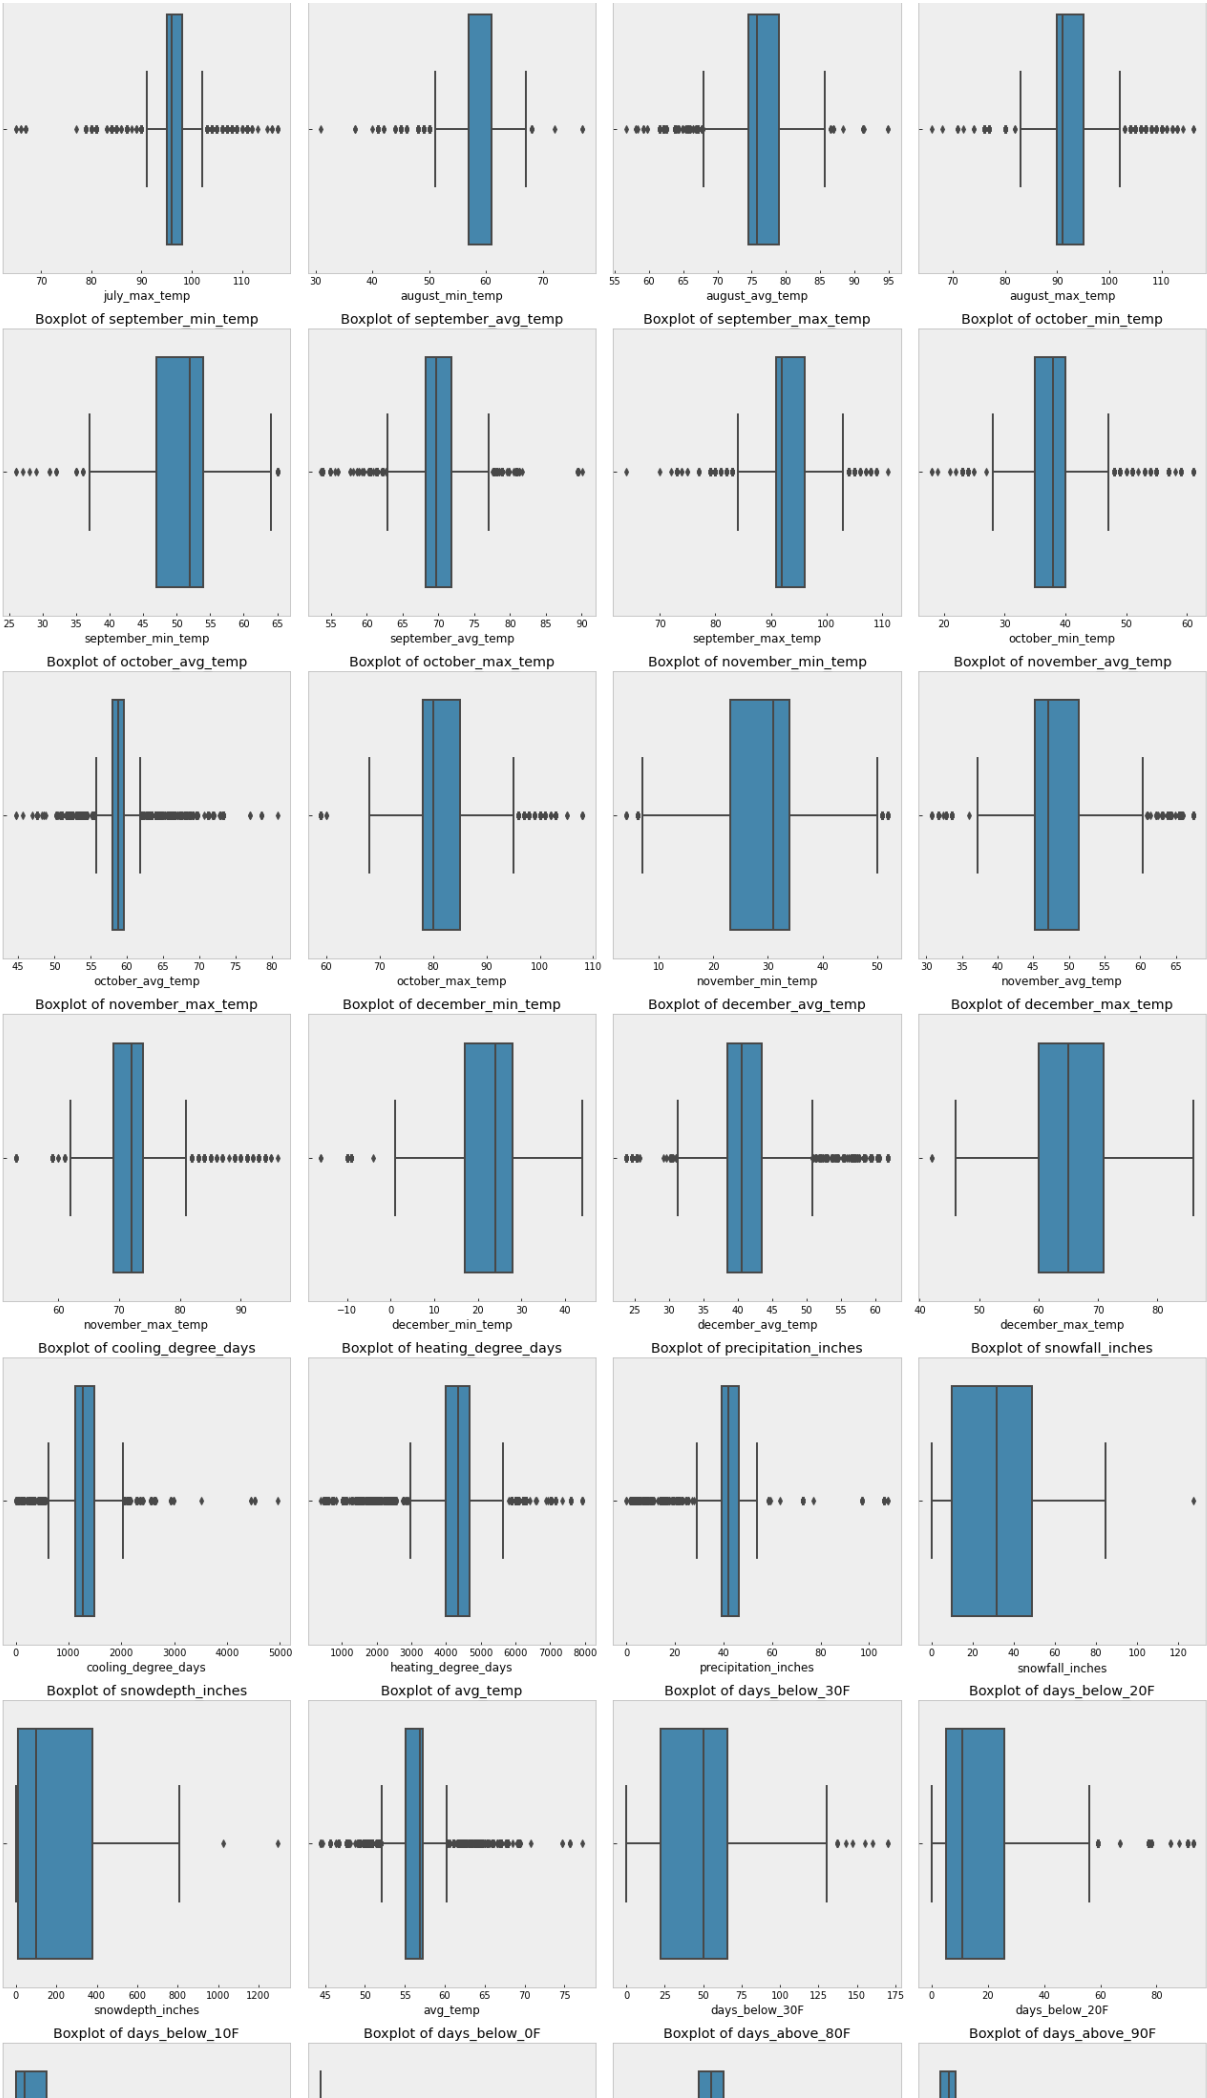
num_features_count = len(Num_features)
num_rows = math.ceil(num_features_count / 4)
fig, axes = plt.subplots(nrows=num_rows, ncols=4, figsize=(18, 5*num_rows))

for i, column in enumerate(Num_features):
    row = i // 4
    col = i % 4
    sns.boxplot(x=energy_df[column], ax=axes[row, col])
    axes[row, col].set_title(f'Boxplot of {column}')
    axes[row, col].set_xlabel(column)

for i in range(num_features_count, num_rows * 4):
    row = i // 4
    col = i % 4
    fig.delaxes(axes[row, col])

plt.tight_layout()
plt.show()
```



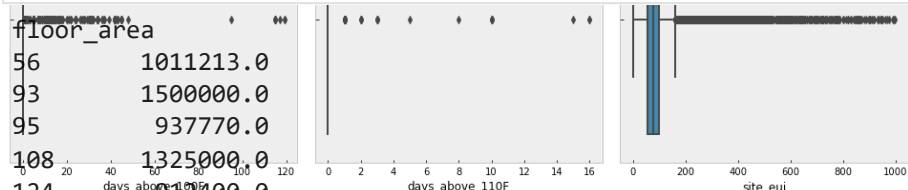
In [31]:

```

""" checking the number of outliers present for each feature using zscore"""
import numpy as np
import scipy.stats

for i in Num_features:
    print(i)
    z = np.abs(scipy.stats.zscore(energy_df[i]))
    outliers = energy_df[z > 3][i]
    print(outliers)

```



Feature	Value
floor_area	56 1011213.0
floor_area	93 150000.0
floor_area	95 937770.0
floor_area	108 1325000.0
floor_area	124 912400.0
...	...
floor_area	73574 1592914.0
floor_area	73654 970647.0
floor_area	73683 962428.0
floor_area	73729 1765970.0
floor_area	73775 2200000.0
year_built	353 0.0
year_built	955 0.0
year_built	2159 0.0
year_built	3415 0.0
year_built	4535 0.0
year_built	5571 0.0
year_built	6931 1789.0
year_built	8411 1789.0
year_built	9348 1829.0
year_built	9907 1789.0
year_built	15123 1600.0
year_built	16936 1836.0
year_built	19948 1600.0
year_built	24302 1649.0
year_built	26477 1827.0
year_built	26790 1836.0
year_built	26876 1600.0
year_built	27251 1649.0
year_built	35003 1827.0
year_built	35472 1600.0
year_built	35896 1649.0
year_built	44412 1827.0
year_built	44853 1836.0
year_built	44951 1600.0
year_built	55459 1827.0
year_built	55686 1800.0
year_built	56141 1836.0
year_built	56260 1600.0
year_built	56441 1811.0
year_built	56882 1649.0
year_built	57581 1833.0
year_built	59640 1800.0
year_built	61605 1799.0
year_built	64737 1841.0
year_built	65735 1818.0
year_built	65907 1756.0
year_built	66037 1800.0

Name: floor_area, Length: 1516, dtype: float64

```
66550    1818.0
66716    1756.0
66851    1800.0
67375    1818.0
67560    1756.0
67713    1800.0
68337    1732.0
68427    1815.0
68592    1756.0
68796    1800.0
Name: year_built, dtype: float64
energy_star_rating
Series([], Name: energy_star_rating, dtype: float64)
ELEVATION
377      958.6
378      958.6
379      958.6
380      958.6
383     1380.7
...
75617    313.0
75618    313.0
75619    313.0
75620    313.0
75621    313.0
Name: ELEVATION, Length: 875, dtype: float64
january_min_temp
358      45
359      45
360      45
361      45
362      45
..
10716   -18
10717   -18
10718   -18
10719   -18
10720   -18
Name: january_min_temp, Length: 1016, dtype: int64
january_avg_temp
358     64.274194
359     64.274194
360     64.274194
361     64.274194
362     64.274194
...
10716    12.258065
10717    12.258065
10718    12.258065
10719    12.258065
10720    12.258065
Name: january_avg_temp, Length: 229, dtype: float64
january_max_temp
358      89
359      89
360      89
361      89
362      89
..
10873     42
10874     42
10875     42
10876     42
10877     42
```

Name: january_max_temp, Length: 1450, dtype: int64

february_min_temp

Series([], Name: february_min_temp, dtype: int64)

february_avg_temp

2280 63.339286

2281 63.339286

2293 63.339286

2294 63.339286

2295 63.339286

2301 63.339286

2316 65.107143

Name: february_avg_temp, dtype: float64

february_max_temp

377 86

378 86

381 85

1061 86

1066 86

..

2374 84

2375 84

2377 84

2378 84

2388 84

Name: february_max_temp, Length: 100, dtype: int64

march_min_temp

2282 52

2283 52

2284 52

2285 52

2286 52

..

12164 -9

12165 -9

12166 -9

12167 -9

12168 -9

Name: march_min_temp, Length: 204, dtype: int64

march_avg_temp

1061 64.548387

1066 64.548387

1067 64.548387

1068 66.096774

1069 64.548387

1070 64.548387

1071 64.548387

1072 64.548387

1073 64.548387

1074 64.548387

1076 64.548387

1077 64.548387

1078 64.548387

1079 64.548387

1087 64.548387

1088 64.548387

1089 64.548387

1090 64.548387

1091 64.548387

1092 64.548387

1093 64.548387

1094 64.548387

1095 64.548387

1110 64.548387

1111 64.548387

```

1112    64.548387
1113    64.548387
1114    64.548387
1115    64.548387
1116    64.548387
1117    64.548387
1121    64.548387
1124    64.548387
1127    64.548387
2280    67.854839
2281    67.854839
2293    67.854839
2294    67.854839
2295    67.854839
2301    67.854839
2316    69.758065

```

Name: march_avg_temp, dtype: float64

march_max_temp

```

381      95
1068     95
2280     94
2281     94
2293     94
2294     94
2295     94
2301     94

```

Name: march_max_temp, dtype: int64

april_min_temp

```

367      49
368      49
369      49
370      49
371      49

```

..

```

2289     52
2290     52
2291     52
2292     52
2404     15

```

Name: april_min_temp, Length: 70, dtype: int64

april_avg_temp

```

377      71.316667
378      71.316667
381      65.316667
382      62.379310
384      62.379310

```

...

```

71585    45.083333
71837    45.083333
72532    45.083333
72533    45.083333
72534    45.083333

```

Name: april_avg_temp, Length: 682, dtype: float64

april_max_temp

```

377      104
378      104
382       99
384       99
385       99

```

...

```

2434      69
2435      69
2436      69
3540      69

```

```
42199      70
Name: april_max_temp, Length: 110, dtype: int64
```

```
may_min_temp
```

```
398      28
```

```
1021     32
```

```
1038     58
```

```
1039     58
```

```
1040     58
```

```
..
```

```
10704    32
```

```
10705    32
```

```
10706    32
```

```
10707    32
```

```
10712    32
```

```
Name: may_min_temp, Length: 80, dtype: int64
```

```
may_avg_temp
```

```
377      80.903226
```

```
378      80.903226
```

```
397      72.322581
```

```
398      51.661290
```

```
399      72.322581
```

```
...
```

```
74808    52.145161
```

```
75061    53.887097
```

```
75754    52.145161
```

```
75755    52.145161
```

```
75756    53.887097
```

```
Name: may_avg_temp, Length: 104, dtype: float64
```

```
may_max_temp
```

```
358      76
```

```
359      76
```

```
360      76
```

```
361      76
```

```
362      76
```

```
..
```

```
74808    79
```

```
75061    80
```

```
75754    79
```

```
75755    79
```

```
75756    80
```

```
Name: may_max_temp, Length: 295, dtype: int64
```

```
june_min_temp
```

```
398      33
```

```
999      39
```

```
1021     34
```

```
1030     36
```

```
1031     36
```

```
..
```

```
3503     38
```

```
3515     38
```

```
3532     31
```

```
3539     34
```

```
42199    40
```

```
Name: june_min_temp, Length: 81, dtype: int64
```

```
june_avg_temp
```

```
0        60.500000
```

```
1        60.500000
```

```
2        60.500000
```

```
3        60.500000
```

```
4        60.500000
```

```
...
```

```
74808    56.233333
```

```
75061    58.433333
```

```
75754    56.233333
```

```
75755    56.233333
75756    58.433333
Name: june_avg_temp, Length: 475, dtype: float64
june_max_temp
358      76
359      76
360      76
361      76
362      76
...
4577    106
4578    106
4579    106
4580    106
4581    106
Name: june_max_temp, Length: 2270, dtype: int64
july_min_temp
398      43
999      46
1012     45
1013     45
1014     45
..
74808    48
75061    48
75754    48
75755    48
75756    48
Name: july_min_temp, Length: 776, dtype: int64
july_avg_temp
0      62.725806
1      62.725806
2      62.725806
3      62.725806
4      62.725806
...
74808    58.758065
75061    60.532258
75754    58.758065
75755    58.758065
75756    60.532258
Name: july_avg_temp, Length: 1464, dtype: float64
july_max_temp
367      81
368      81
369      81
370      81
371      81
..
74808    81
75061    83
75754    81
75755    81
75756    83
Name: july_max_temp, Length: 1210, dtype: int64
august_min_temp
377      77
378      77
398      44
999      45
1021     45
..
68167    41
68168    41
```

```
68169    41
68170    41
68171    41
Name: august_min_temp, Length: 1575, dtype: int64
august_avg_temp
0      62.161290
1      62.161290
2      62.161290
3      62.161290
4      62.161290
...
74806    61.612903
74807    61.612903
74808    61.612903
75754    61.612903
75755    61.612903
Name: august_avg_temp, Length: 1465, dtype: float64
august_max_temp
377     116
378     116
379     105
380     105
382     108
...
5614      76
5615      76
5616      76
5617      76
33196     80
Name: august_max_temp, Length: 2137, dtype: int64
september_min_temp
367      65
368      65
369      65
370      65
371      65
372      65
373      65
374      65
375      65
377      65
378      65
1119     27
1223     31
2244     32
2245     32
2249     32
2250     32
2258     36
2259     36
2266     36
2267     36
2270     29
2272     26
2273     26
2282     64
2283     64
2284     64
2285     64
2286     64
2287     64
2288     64
2289     64
2290     64
```

2291	64
2292	64
2316	65
2404	31
2431	35
2434	35
2435	35
2436	35
3532	28
3539	36
65627	37
65628	37
69347	37
69354	37
70058	37
70742	37
70743	37
71274	37
71583	37
71584	37
71585	37
71837	37
72532	37
72533	37
72534	37

Name: september_min_temp, dtype: int64

september_avg_temp

377	89.550000
378	89.550000
381	80.500000
388	80.950000
389	80.950000

...

74808	53.783333
75061	55.931034
75754	53.783333
75755	53.783333
75756	55.931034

Name: september_avg_temp, Length: 71, dtype: float64

september_max_temp

377	109
378	109
381	108
1016	64
1068	108
2277	74
2278	74
2279	70
2280	109
2281	109
2293	109
2294	109
2295	109
2301	109
2306	108
2307	108
2316	111
2319	108
2320	108
2365	108
2368	108
2369	108
2370	108
2393	77

2394	77
2395	77
2399	77
2401	77
2402	77
2403	77
2405	77
2406	77
3540	72
72538	73
72544	73
72545	73
73255	73
73946	73
73947	73
74482	73
74796	73
74797	73
74798	73
74799	73
74800	73
74801	73
74802	73
74803	73
74804	73
74805	73
74806	73
74807	73
74808	73
75061	75
75754	73
75755	73
75756	75

Name: september_max_temp, dtype: int64

october_min_temp

367	54
368	54
369	54
370	54
371	54
..	
3459	55
3460	55
3461	55
3462	55
52611	21

Name: october_min_temp, Length: 1014, dtype: int64

october_avg_temp

358	69.774194
359	69.774194
360	69.774194
361	69.774194
362	69.774194
...	
74808	47.661290
75061	48.532258
75754	47.661290
75755	47.661290
75756	48.532258

Name: october_avg_temp, Length: 2196, dtype: float64

october_max_temp

377	108
378	108
381	105

```
382      108
384      108
...
74808    59
75061    60
75754    59
75755    59
75756    60
Name: october_max_temp, Length: 761, dtype: int64
november_min_temp
1038     51
1039     51
1040     51
1041     51
1042     51
..
10772     6
10773     6
10774     6
10775     6
10776     6
Name: november_min_temp, Length: 307, dtype: int64
november_avg_temp
358     63.016667
359     63.016667
360     63.016667
361     63.016667
362     63.016667
...
10716    31.716667
10717    31.716667
10718    31.716667
10719    31.716667
10720    31.716667
Name: november_avg_temp, Length: 423, dtype: float64
november_max_temp
358     90
359     90
360     90
361     90
362     90
..
71585    53
71837    53
72532    53
72533    53
72534    53
Name: november_max_temp, Length: 222, dtype: int64
december_min_temp
2272    -16
2273    -16
12177   -10
12178   -10
12179   -10
..
14784    -9
14785    -9
14786    -9
14787    -9
14788   -10
Name: december_min_temp, Length: 2614, dtype: int64
december_avg_temp
1038    59.387097
1039    59.387097
```

```

1040    59.387097
1041    59.387097
1042    59.387097
...
14719   23.790323
14720   23.790323
14721   23.790323
14722   23.790323
14788   23.790323
Name: december_avg_temp, Length: 401, dtype: float64
december_max_temp
1061     85
1066     85
1067     85
1069     85
1070     85
..
74808    46
75061    42
75754    46
75755    46
75756    42
Name: december_max_temp, Length: 134, dtype: int64
cooling_degree_days
377    4453
378    4453
388    2579
389    2579
393    2579
...
3533     4
3536     4
3537     4
3538     4
3540     1
Name: cooling_degree_days, Length: 62, dtype: int64
heating_degree_days
358    1125
359    1125
360    1125
361    1125
362    1125
...
10717   7580
10718   7580
10719   7580
10720   7580
42465   6933
Name: heating_degree_days, Length: 2599, dtype: int64
precipitation_inches
358    10.43
359    10.43
360    10.43
361    10.43
362    10.43
...
74808   106.32
75061   107.69
75754   106.32
75755   106.32
75756   107.69
Name: precipitation_inches, Length: 2453, dtype: float64
snowfall_inches
398     84.8

```

```
2404      127.3
Name: snowfall_inches, dtype: float64
snowdepth_inches
2270      1023
2404      1292
52611     807
Name: snowdepth_inches, dtype: int64
avg_temp
358      64.251366
359      64.251366
360      64.251366
361      64.251366
362      64.251366
...
74808     47.911202
75061     49.127397
75754     47.911202
75755     47.911202
75756     49.127397
Name: avg_temp, Length: 1858, dtype: float64
days_below_30F
1119      147
2270      155
2272      170
2273      170
2404      160
3532      143
10696     137
10697     137
10704     137
10705     137
10706     137
10707     137
10712     137
Name: days_below_30F, dtype: int64
days_below_20F
2272      85
2273      85
10489     78
10490     78
10491     78
..
10718     91
10719     91
10720     91
42465     67
52611     67
Name: days_below_20F, Length: 236, dtype: int64
days_below_10F
2272      29
2273      29
10489     45
10490     45
10491     45
..
68167     30
68168     30
68169     30
68170     30
68171     30
Name: days_below_10F, Length: 2536, dtype: int64
days_below_0F
2272      22
2273      22
```

```

10489    25
10490    25
10491    25
...
14720    12
14721    12
14722    12
14788    12
52611    12
Name: days_below_0F, Length: 1980, dtype: int64
days_above_80F
377      246
378      246
379      176
380      176
381      162
...
74806     6
74807     6
74808     6
75754     6
75755     6
Name: days_above_80F, Length: 270, dtype: int64
days_above_90F
377      182
378      182
379      113
380      113
381       54
...
4581      52
69341     52
69342     52
69343     52
69344     52
Name: days_above_90F, Length: 1285, dtype: int64
days_above_100F
377      119
378      119
379       12
380       12
382       30
...
4577      10
4578      10
4579      10
4580      10
4581      10
Name: days_above_100F, Length: 1302, dtype: int64
days_above_110F
377       16
378       16
394        1
395        1
396        1
...
2301      10
2316      15
3531       2
3534       2
3535       2
Name: days_above_110F, Length: 61, dtype: int64
site_eui
13      608.839519

```

```

24      287.863448
26      264.068722
113     275.977289
144     285.862933

```

```

...
75376    268.596672
75424    364.958302
75456    268.380928
75520    275.649545
75755    592.022750

```

```
Name: site_eui, Length: 1216, dtype: float64
```

Dropping Columns Based on Box Plot Observations

Upon visual inspection using box plots, it's evident that the columns 'days_above_110F' and 'days_above_100F' and 'days_below_0F' predominantly contain zero values for most data points.

Therefore, considering the lack of variability and information in these columns, I have decided to drop both 'days_above_110F' and 'days_above_100F' and 'days_below_0F' from the dataset.

```
In [32]: energy_df.drop(['days_above_110F', 'days_above_100F', 'days_below_0F'], axis=1, inplace=True)
```

```
In [33]: energy_df.shape
```

```
Out[33]: (75757, 56)
```

During the analysis, numerous outliers were detected across all features. However, the presence of a large number of outliers does not necessarily indicate their irrelevance.

For this analysis, I have decided not to remove outliers as their presence might hold valuable information or characteristics within the dataset. Retaining outliers can contribute to a more comprehensive understanding of the data and potentially enhance the performance of the models.

```
In [ ]:
```

```
In [34]: import pandas as pd

# Assuming energy_df_id_remove_2 is your DataFrame and features is a List of column names
features = energy_df.columns.tolist()

Num_features = [feature for feature in features if
                 energy_df[feature].dtype != 'object' and
                 feature != 'Year_Factor']

Cat_features = [feature for feature in features if
                energy_df[feature].dtype == object or feature == 'Year_Factor' ]

print(Num_features)
print("\n")

print(Cat_features)
```

```

['floor_area', 'year_built', 'energy_star_rating', 'ELEVATION', 'january_min_temp',
 'january_avg_temp', 'january_max_temp', 'february_min_temp', 'february_avg_temp', 'february_max_temp',
 'march_min_temp', 'march_avg_temp', 'march_max_temp', 'april_min_temp', 'april_avg_temp', 'april_max_temp',
 'may_min_temp', 'may_avg_temp', 'may_max_temp', 'june_min_temp', 'june_avg_temp', 'june_max_temp',
 'july_min_temp', 'july_avg_temp', 'july_max_temp', 'august_min_temp', 'august_avg_temp', 'august_max_temp',
 'september_min_temp', 'september_avg_temp', 'september_max_temp', 'october_min_temp', 'october_avg_temp',
 'october_max_temp', 'november_min_temp', 'november_avg_temp', 'november_max_temp', 'december_min_temp',
 'december_avg_temp', 'december_max_temp', 'Year_Factor']

```

```
mp', 'april_avg_temp', 'april_max_temp', 'may_min_temp', 'may_avg_temp', 'may_max_tem  
p', 'june_min_temp', 'june_avg_temp', 'june_max_temp', 'july_min_temp', 'july_avg_tem  
p', 'july_max_temp', 'august_min_temp', 'august_avg_temp', 'august_max_temp', 'septem  
ber_min_temp', 'september_avg_temp', 'september_max_temp', 'october_min_temp', 'octob  
er_avg_temp', 'october_max_temp', 'november_min_temp', 'november_avg_temp', 'november  
_max_temp', 'december_min_temp', 'december_avg_temp', 'december_max_temp', 'cooling_d  
egree_days', 'heating_degree_days', 'precipitation_inches', 'snowfall_inches', 'snowd  
epth_inches', 'avg_temp', 'days_below_30F', 'days_below_20F', 'days_below_10F', 'days  
_above_80F', 'days_above_90F', 'site_eui']
```

```
['Year_Factor', 'State_Factor', 'building_class', 'facility_type']
```

In []:


Correlation matrix to check highly correlated features

In [35]:

```
corr = energy_df[Num_features].corr()  
plt.subplots(figsize=(60, 60))  
sns.heatmap(corr, linewidths=.5, vmin=0, vmax=1, square=True)
```

Out[35]:

<AxesSubplot:>



I will remove highly_correlated features(threshold = 0.9)

In [36]:

```
# Assuming 'corr' is the correlation matrix calculated from Num_features

threshold = 0.9 # Set your desired threshold for correlation

# Create a mask to focus only on the upper triangle of the correlation matrix (to av
mask = np.triu(np.ones_like(corr, dtype=bool), k=1)

# Find columns with correlation above the threshold
highly_correlated = set()
for i in range(len(corr.columns)):
    for j in range(i+1, len(corr.columns)):
        if mask[i, j] and abs(corr.iloc[i, j]) > threshold:
            col_i = corr.columns[i]
            col_j = corr.columns[j]
            highly_correlated.add(col_i)
            highly_correlated.add(col_j)

#
highly_correlated.add(col_j)

print("Columns highly correlated:", highly_correlated)
```

Columns highly correlated: {'january_min_temp', 'february_avg_temp', 'december_min_te
mp', 'days_below_20F', 'january_avg_temp', 'august_avg_temp', 'july_avg_temp'}


```
In [37]: len(highly_correlated)
```

```
Out[37]: 7
```

```
In [38]: """There are 6 numerical features are highly correlated with other data"""
        """So I am removing those 6 columns """

        energy_df.drop(highly_correlated, inplace=True, axis=1)
```

```
In [39]: energy_df.shape
```

```
Out[39]: (75757, 49)
```

Numerical features and Categorical features

```
In [40]: import pandas as pd

        # Assuming energy_df_id_remove_2 is your DataFrame and features is a List of column names
        features = energy_df.columns.tolist()

        Num_features = [feature for feature in features if
                        energy_df[feature].dtype != 'object' and
                        feature != 'Year_Factor' and feature != 'site_eui']
        Cat_features = [feature for feature in features if
                        energy_df[feature].dtype == object or feature == 'Year_Factor' ]

        print(Num_features)
        print("\n")

        print(Cat_features)

        ['floor_area', 'year_built', 'energy_star_rating', 'ELEVATION', 'january_max_temp',
        'february_min_temp', 'february_max_temp', 'march_min_temp', 'march_avg_temp', 'march_
        max_temp', 'april_min_temp', 'april_avg_temp', 'april_max_temp', 'may_min_temp', 'may
        _avg_temp', 'may_max_temp', 'june_min_temp', 'june_avg_temp', 'june_max_temp', 'july_
        min_temp', 'july_max_temp', 'august_min_temp', 'august_max_temp', 'september_min_tem
        p', 'september_avg_temp', 'september_max_temp', 'october_min_temp', 'october_avg_tem
        p', 'october_max_temp', 'november_min_temp', 'november_avg_temp', 'november_max_tem
        p', 'december_avg_temp', 'december_max_temp', 'cooling_degree_days', 'heating_degree_
        days', 'precipitation_inches', 'snowfall_inches', 'snowdepth_inches', 'avg_temp', 'da
        ys_below_30F', 'days_below_10F', 'days_above_80F', 'days_above_90F']

        ['Year_Factor', 'State_Factor', 'building_class', 'facility_type']
```

```
In [ ]:
```

Splitting the Data into Train and Test Data

```
In [41]: X = energy_df.drop(['site_eui'],axis=1)
        y = energy_df['site_eui']
        X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=42, test_size=0.2)
```

```
In [42]: X_train.shape, X_test.shape, y_train.shape, y_test.shape, X.shape, y.shape
```

```
Out[42]: ((60605, 48), (15152, 48), (60605,), (15152,), (75757, 48), (75757,))
```

Plot the decision tree

Data Preprocessing for Decision Tree regressor

For the decision tree Regressor model:

1. **Normalization:** Decision trees do not require normalization as they are not sensitive to the scale of numerical features. Hence, normalization is not necessary for this model.
2. **One-Hot Encoding:** Categorical features need to be one-hot encoded as decision trees typically require categorical variables to be converted into a numerical format for processing.

Therefore, before fitting the data into the decision tree model:

- Apply one-hot encoding to categorical features.
- No need to perform normalization on numerical features.

```
In [43]: full_pipeline_1 = ColumnTransformer([
#      ('StandardScale', StandardScaler(), Num_features),
      ('onehot', OneHotEncoder(), Cat_features)
],remainder='passthrough')
```

```
In [44]: pipe1= Pipeline(steps=[('full_pipeline_1',full_pipeline_1),
                                ])
```

```
In [ ]:
```

```
In [45]: X_train_transformed = pipe1.fit_transform(X_train)
```

```
In [46]: X_test_transformed = pipe1.transform(X_test)
```

```
In [47]: # X_train_transformed
```

```
In [48]: from sklearn.tree import DecisionTreeClassifier, DecisionTreeRegressor, plot_tree
```

```
In [ ]:
```

```
In [49]: from sklearn.tree import DecisionTreeRegressor, plot_tree # Add plot_tree import
import matplotlib.pyplot as plt

def train_using_entropy(X_train, y_train):
```

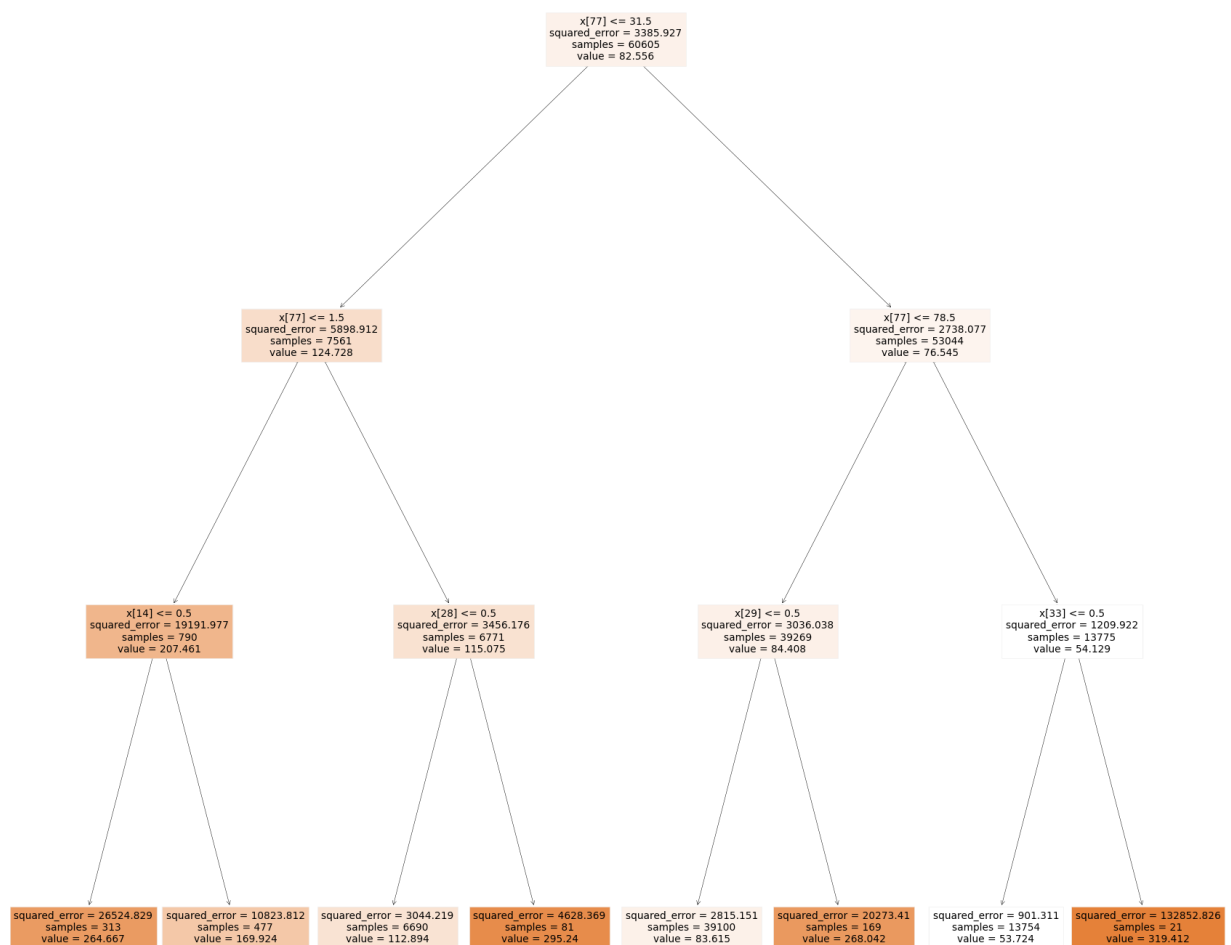
```
# Decision tree with entropy
clf_entropy = DecisionTreeRegressor(max_depth=3, criterion='squared_error')

# Performing training
clf_entropy.fit(X_train, y_train)

plt.figure(figsize=(30, 30))
plot_tree(clf_entropy, filled=True) # Use plot_tree from sklearn.tree
plt.show()

return clf_entropy
```

```
clf_object = train_using_entropy(X_train_transformed, y_train)
```



Here for this Decision tree i have used depth as 3 just to visualize the tree. But for proper data prediction I should choose much higher depth value otherwise prediction will be not properly correct

In [50]:

```
def prediction(X_test, clf_object):

    y_pred = clf_object.predict(X_test)
    mse = mean_squared_error(y_test, y_pred)
    print("Mean Squared Error:", mse)
    print("Predicted values:")
```

```
print(y_pred)
return y_pred
```

```
In [51]: y_pred_entropy = prediction(X_test_transformed, clf_object)
```

Mean Squared Error: 2764.9598676970622

Predicted values:

```
[83.6145246  53.72382701 83.6145246 ... 83.6145246  83.6145246
 53.72382701]
```

```
In [ ]:
```

Plot the PCA

Principal Component Analysis (PCA) is a dimensionality reduction technique used to reduce the number of features (variables) while retaining the most important information or patterns present in the original dataset.

Normalizing the data and handling categorical features are important steps before applying PCA.

```
In [52]: full_pipeline_2 = ColumnTransformer([
    ('StandardScale', StandardScaler(), Num_features),
    ('onehot', OneHotEncoder(), Cat_features),

])
```

```
In [53]: pipe2= Pipeline(steps=[('full_pipeline_2',full_pipeline_2),
    # ('pca' , PCA(n_components=0.90, svd_solver='full'))

])
```

```
In [54]: pipe2
```

```
Out[54]: Pipeline
└─ full_pipeline_2: ColumnTransformer
    ├── StandardScale
    │   └─ StandardScaler
    └─ onehot
        └─ OneHotEncoder
```

```
In [55]: X_train_transformed2 = pipe2.fit_transform(X_train)
X_test_transformed2 = pipe2.transform(X_test)
```

```
In [ ]:
```

```
In [56]: pca_model = PCA(n_components=0.90, svd_solver='full')
X_train_pca = pca_model.fit_transform(X_train_transformed2)
X_test_pca = pca_model.transform(X_test_transformed2)
```

```
In [57]: X_train_pca.shape, X_test_pca.shape
```

```
Out[57]: ((60605, 12), (15152, 12))
```

In your case, if you've conducted PCA and it resulted in reducing the number of features to 12, it means that these 12 components retain the most relevant information from the original dataset as I am using 90% data relevancy.

```
In [58]: pca_model.explained_variance_ratio_
```

```
Out[58]: array([0.26600548, 0.19917598, 0.12599467, 0.07552799, 0.06216106,
        0.04343502, 0.03120888, 0.02515717, 0.02320775, 0.02275822,
        0.01984301, 0.01830348])
```

```
In [59]: # %matplotlib inline
import matplotlib.pyplot as plt
plt.rcParams["figure.figsize"] = (12,6)


fig, ax = plt.subplots()
xi = np.arange(1, 13, step=1) # X_train_pca.shape =(60605, 12) thats why 1 am us
y = np.cumsum(pca_model.explained_variance_ratio_)

plt.ylim(0.0,1.1)
plt.plot(xi, y, marker='o', linestyle='--', color='b')

plt.xlabel('Number of Components')
plt.xticks(np.arange(0, 13, step=1)) #change from 0-based array index to 1-based hum
plt.ylabel('Cumulative variance (%)')
plt.title('The number of components needed to explain variance')

plt.axhline(y=0.90, color='r', linestyle='-')
plt.text(0.5, 0.85, '90% cut-off threshold', color = 'red', fontsize=16)


ax.grid(axis='x')
plt.show()
```



```
In [60]: import matplotlib.pyplot as plt
from sklearn.decomposition import PCA
import pandas as pd
from sklearn.preprocessing import StandardScaler

# pca_model.plot(figsize=(10,8))
# plt.show()


# Plotting the explained variance ratio
plt.figure(figsize=(8, 6))
plt.bar(range(1, len(pca_model.explained_variance_ratio_) + 1), pca_model.explained_
plt.xlabel('Principal Component')
plt.ylabel('Explained Variance Ratio')
plt.title('Explained Variance Ratio of Principal Components')
plt.show()
```



```
In [61]: explained_variance = pca_model.explained_variance_ratio_
singular_values = pca_model.singular_values_
```

```
In [62]: x = np.arange(1, len(explained_variance)+1)
plt.bar(range(1, len(pca_model.explained_variance_ratio_) + 1), pca_model.explained_

plt.plot(x, explained_variance)
plt.ylabel('Share of Variance Explained')
plt.title("PCA explained variance plot")
plt.xlabel("Components")
plt.show()
```



```
In [63]: for i in range(0, 12):
          print(f"Component {i:>2} accounts for {explained_variance[i]*100:>2.2f}% of vari
```

```
Component 0 accounts for 26.60% of variance
Component 1 accounts for 19.92% of variance
Component 2 accounts for 12.60% of variance
Component 3 accounts for 7.55% of variance
Component 4 accounts for 6.22% of variance
Component 5 accounts for 4.34% of variance
Component 6 accounts for 3.12% of variance
Component 7 accounts for 2.52% of variance
Component 8 accounts for 2.32% of variance
Component 9 accounts for 2.28% of variance
Component 10 accounts for 1.98% of variance
Component 11 accounts for 1.83% of variance
```

Choosing the Optimal Number of Principal Components

When working with Principal Component Analysis (PCA) for dimensionality reduction, deciding the number of components is crucial. Here are some methods to determine the optimal number of principal components:

1. Examining the Knee in Explained Variance Plot:

- In our dataset, the explained variance plot exhibits a noticeable "knee" around 4-6 principal components. This knee point can be indicative of the optimal number of components.

2. Keeping Components Explaining Significant Variance:

- Another approach involves retaining components that account for more than 1% of the variance in the dataset. For our data, this threshold occurs after 11 components.

3. Retaining Components with Cumulative Explained Variance:

- Considering the cumulative explained variance, it's beneficial to retain principal components that collectively cover a substantial portion of the total variance. For instance, keeping components that contribute to approximately 80% of the explained variance in the dataset. In our case, this would encompass the first 7 components.

These methods assist in striking a balance between computational efficiency and model performance by selecting an appropriate number of principal components for dimensionality reduction.

From this PCA Components plot we can understand that $n_{\text{components}} = 12$ will capture 90% of feature knowledge

LinearRegression model is trained after PCA

```
In [64]: # Creating the final pipeline with linear regression model
pipeline = Pipeline(steps=[('regressor', LinearRegression())])
pipeline.fit(X_train_pca, y_train)

# Evaluating the model
train_score = pipeline.score(X_train_pca, y_train)
test_score = pipeline.score(X_test_pca, y_test)

print(f"Training R^2 score: {train_score:.4f}")
print(f"Testing R^2 score: {test_score:.4f}")
```

Training R² score: 0.1669

Testing R² score: 0.1703

```
In [65]: from sklearn.metrics import mean_squared_error
import numpy as np

y_pred = pipeline.predict(X_test_pca)

# Calculate Mean Squared Error
mse = mean_squared_error(y_test, y_pred)

print(f" Mean Squared Error (MSE): {mse:.4f}")
```

Mean Squared Error (MSE): 2841.2050

So here I got R² score: 0.1703 and Mean Squared Error is 2841 which is very poor performance with simple linear regression.

PLOT THE DATA WITH 2 PCA COMPONENTS

```
In [66]: pca_2 = PCA(n_components=2, whiten=True)
#fit the model to our data and extract the results
X_pca_2 = pca_2.fit_transform(X_train_transformed2)
```

```
In [67]: df_pca_plot = pd.DataFrame(data = X_pca_2,
                                columns = ["Component 1",
                                           "Component 2"])
```

```
In [68]: df_pca_plot["Component 1"]
```


```
Out[68]: 0      -0.972568
1      -0.884896
2      -0.059776
3      -0.045234
4      -0.697619
```



```
...
60600    -0.746157
60601    -0.460480
60602    -0.012080
60603     1.946435
60604    -0.188181
Name: Component 1, Length: 60605, dtype: float64
```

In [69]:

```
#plot the resulting data from two dimensions
plot = plt.scatter(df_pca_plot["Component 1"], df_pca_plot["Component 2"], c=y_train)
plt.legend(handles=plot.legend_elements()[0],)
plt.xlabel("Principal Component 1") # Naming the x-axis
plt.ylabel("Principal Component 2") # Naming the y-axis
plt.show()
```



Plot t-SNE on the PCA-transformed data

Dimensionality reduction using PCA (Principal Component Analysis) followed by t-SNE (t-distributed Stochastic Neighbor Embedding) for visualizing the dataset in a 2D space.


In [70]:

```
from sklearn.decomposition import PCA
from sklearn.manifold import TSNE
import matplotlib.pyplot as plt
import seaborn as sns

# Apply PCA
pca_t = PCA(n_components=2) # Choose the number of components
X_pca = pca_t.fit_transform(X_train_transformed2)

# Apply t-SNE on the PCA-transformed data
tsne = TSNE(n_components=2, random_state=42)
X_tsne = tsne.fit_transform(X_pca)

# Plotting the data in 2D
plt.figure(figsize=(8, 6))
sns.scatterplot(x=X_tsne[:,0], y=X_tsne[:,1], palette='viridis')
plt.title('t-SNE Visualization of Data after PCA')
plt.xlabel('t-SNE Component 1')
plt.ylabel('t-SNE Component 2')
plt.show()
```



It shows different similar data clusters but as I consider whole training dataset that's why it's not able to show that properly. But if I use a small part of dataset it will show the clusters properly.

In []:

In []:

Building Pipeline

Using Scikit-Learn's Pipeline for Machine Learning Workflows

Scikit-Learn provides a powerful tool called `Pipeline` that allows you to chain multiple steps together for a machine learning workflow. These steps can include preprocessing, feature selection, and model building.

Benefits of Using Pipelines:

1. **Simplified Workflow:**

- Pipelines allow you to combine several data processing steps into a single object, making it easier to manage and reproduce the workflow.

2. **Preventing Data Leakage:**

- Pipelines help in avoiding data leakage by ensuring that preprocessing steps (e.g., scaling, imputation) are applied consistently to training and testing data.

3. **Cross-Validation Handling:**


- It's simpler to perform cross-validation with a pipeline since transformations occur within each fold, preventing data leakage and ensuring a more accurate estimation of model performance.

```
In [71]: energy_df_pipe = pd.read_csv("C:\\Users\\praba\\Desktop\\uca1\\M1\\ML\\final project
energy_df_pipe.tail(2)
```

```
Out[71]:
```

	Year_Factor	State_Factor	building_class	facility_type	floor_area	year_built	energy_star
75755	6	State_11	Commercial	Commercial_Other	53575.0	2012.0	
75756	6	State_11	Residential	2to4_Unit_Building	23888.0	1974.0	

2 rows × 64 columns



```
In [72]: X_1 = energy_df_pipe.drop(['site_eui'],axis=1)
y_1 = energy_df_pipe['site_eui']
X_train_pipe, X_test_pipe, y_train_pipe, y_test_pipe = train_test_split(X_1, y_1 , r
```

```
In [73]: X_train_pipe.shape, y_test_pipe.shape
```

```
Out[73]: ((60605, 63), (15152,))
```

These drop_highly_correlated_columns I found after using coorelation matrix and setting threshold =90

```
In [74]: drop_highly_correlated_columns =['january_avg_temp', 'december_min_temp', 'august_av
```

```
In [75]: len(drop_highly_correlated_columns)
```

```
Out[75]: 6
```

```
In [ ]:
```

If a feature contains a vast majority of missing values (for instance, more than 40,000 null values in this case), one common approach is to consider dropping those features from the dataset.

These 4 features contain more tahn 40000 Null Values.

```
In [76]: drop_high_null_valued_column = ['direction_max_wind_speed','direction_peak_wind_spee
```

all ids are unique so not relavent

```
In [77]: drop_unique_id =['id']
```

it's evident that the columns 'days_above_110F' and 'days_above_100F' and 'days_below_0F' predominantly contain zero values for most data points. so i am removing

```
In [78]: drop_ulrelated = ['days_above_110F', 'days_above_100F', 'days_below_0F']
```

Why I am removing those features are explained in data preprocessing step

```
In [79]: drop_features_1= drop_highly_correlated_columns+drop_high_null_valued_column+drop_un
print(drop_features_1)
print(len(drop_features_1))
```

```
['january_avg_temp', 'december_min_temp', 'august_avg_temp', 'july_avg_temp', 'januar
y_min_temp', 'days_below_20F', 'direction_max_wind_speed', 'direction_peak_wind_spee
d', 'max_wind_speed', 'days_with_fog', 'id', 'days_above_110F', 'days_above_100F', 'd
ays_below_0F']
14
```

```
In [80]: import pandas as pd

# Assuming energy_df_id_remove_2 is your DataFrame and features is a list of column
features = X_1.columns.tolist()

Numerical_features_1 = [feature for feature in features if
                        X_1[feature].dtype != 'object' and
                        feature != 'Year_Factor']
categorical_features_1 = [feature for feature in features if
                          X_1[feature].dtype == object or feature == 'Year_Factor' ]

print(len(Numerical_features_1))
# print("\n")

print(len(categorical_features_1))
```

59

4

```
In [81]: # Numerical_features_1
```

```
In [82]: # Numerical_features
```

```
In [83]: categorical_features_1=['Year_Factor', 'State_Factor', 'building_class', 'facility_t
```

```
In [84]: common_elements_1 = list(set(Numerical_features_1) & set(drop_features_1))
print(common_elements_1)
```

```
['january_min_temp', 'id', 'days_with_fog', 'july_avg_temp', 'december_min_temp', 'da
ys_below_0F', 'days_above_110F', 'days_below_20F', 'january_avg_temp', 'days_above_10
0F', 'direction_peak_wind_speed', 'direction_max_wind_speed', 'august_avg_temp', 'max
_wind_speed']
```

```
In [85]: # Remove elements from list1 that are present in list2
Numerical_features_2 = [x for x in Numerical_features_1 if x not in common_elements_
print(len(Numerical_features_2))
```

45

```
In [86]: common_elements_2 = list(set(Numerical_features_2) & set(drop_features_1))
print(common_elements_2)
```

```
[]
```

```
In [87]: X_train_pipe.shape, X_test_pipe.shape, y_train_pipe.shape, y_test_pipe.shape
```

```
Out[87]: ((60605, 63), (15152, 63), (60605,), (15152,))
```

```
In [88]: common_elements = list(set(Numerical_features_2) & set(drop_features_1))
print(common_elements)
```

```
[]
```

```
In [89]: """This is the pipeline To drop specific columns """
drop_transformer = ColumnTransformer(transformers=[('drop_columns','drop',drop_featu
```

```
In [ ]:
```

```
In [90]: # drop_transformer = ColumnTransformer(transformers=[('drop_columns','drop',drop_fea

# Creating pipelines for numerical and categorical features
numerical_pipeline = Pipeline(steps=[
#     ('outlier_removal', remove_outliers1(remove_outliers)),
    ('imputer', SimpleImputer(strategy='mean')) ,
    ('stdscaler', StandardScaler())
])

categorical_pipeline = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='most_frequent')), # Filling missing values
    ('onehot', OneHotEncoder(handle_unknown='ignore')) # One-hot encoding categoric
])

# Creating a ColumnTransformer to apply the appropriate pipeline to each type of fea
col_transformer = ColumnTransformer(transformers=[
    ('drop_columns','drop',drop_features_1),
    ('numerical', numerical_pipeline, Numerical_features_2),
    ('categorical', categorical_pipeline, categorical_features_1),
#     ('scale', StandardScaler())
],remainder='drop')

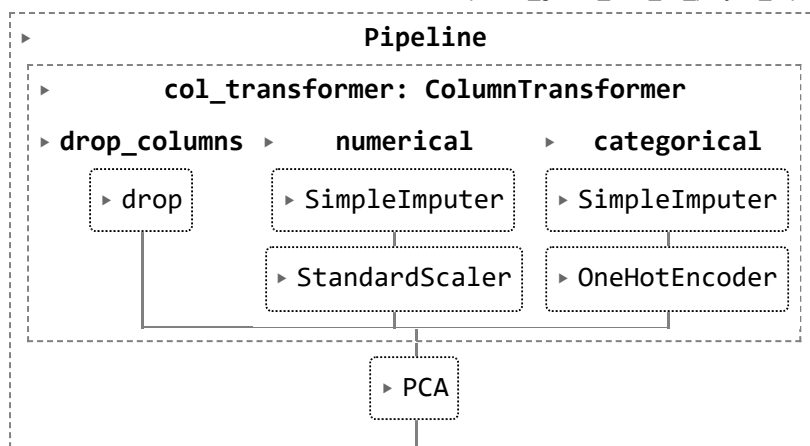
pipe = Pipeline(steps=[

#     ('drop_transformer',drop_transformer),
#         ('col_transformer',col_transformer),

    ('pca',PCA()),
])
```

```
In [91]: pipe
```

Out[91]:



Optimizing Model Performance with Hyperparameter Tuning

In [92]:

```

# Initialize the estimators
clf1 = RandomForestRegressor()
clf2 = Lasso()

clf3 = LinearRegression()
clf4 = Ridge()
clf5 = ElasticNet()

clf6 = DecisionTreeRegressor()

```

In [93]:

```

# Initiaze the hyperparameters for each dictionary

#hyperparameters for RandomForestRegressor
param1 = {}
param1['regressor__n_estimators'] = [10,20]
param1['regressor__max_depth'] = [10,15]
param1['regressor'] = [clf1]

#hyperparameters for Lasso

param2 = {}
param2['regressor__alpha'] = [0.1, 1, 10]
param2['regressor'] = [clf2]

#hyperparameters for LinearRegression

param3 = {}
param3['regressor'] = [clf3]

#hyperparameters for Ridge

param4 = {}
param4['regressor__alpha'] = [0.1, 1]
param4['regressor'] = [clf4]

#hyperparameters for ELasticNet

param5 = {}
param5['regressor__alpha'] = [0.1, 1]
param5['regressor__l1_ratio'] = [0.2, 0.5, 0.7]
param5['regressor'] = [clf5]

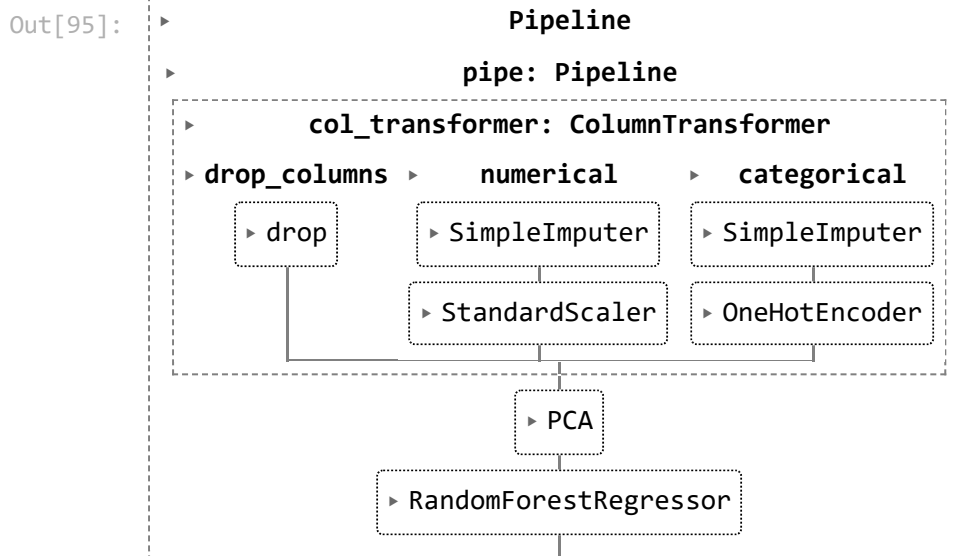
#hyperparameters for DecisionTreeRegressor

```

```
param6 = {}
param6['regressor__max_depth'] = [3, 6]
param6['regressor'] = [clf6]
```

```
In [94]: pipeline = Pipeline(steps=[('pipe',pipe),
                                   ('regressor', clf1)])
# params = [param1, param2, param3, param4, param5,param6]
params = [param1, param2, param3]
```

```
In [95]: pipeline
```



```
In [ ]:
```

```
In [96]: # Gridsearchcv is used
# %%time

gs = GridSearchCV(pipeline, params, cv=3, n_jobs=-1, scoring='neg_root_mean_squared_
```

```
In [97]: # X_train_pipe
```

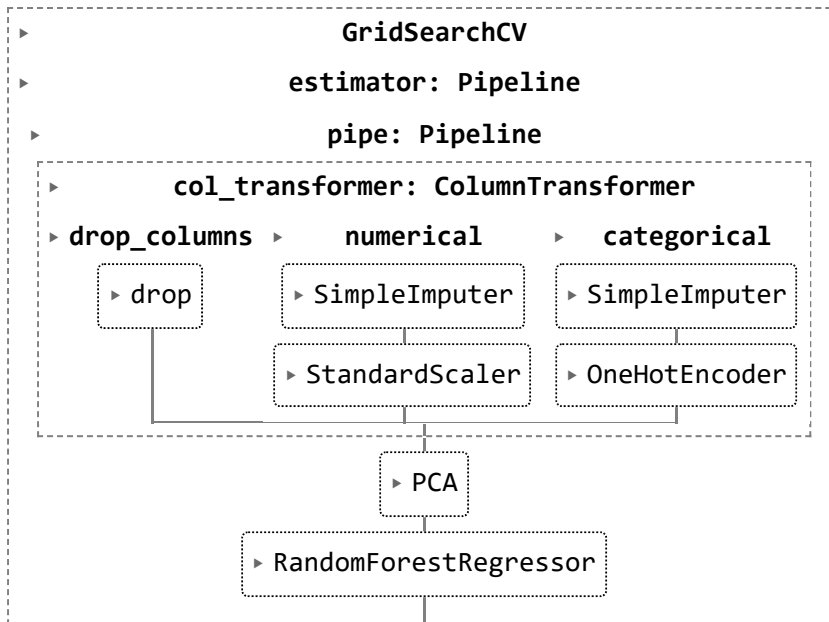
```
In [98]: %%time

gs_fit = gs.fit(X_train_pipe, y_train_pipe)

Fitting 3 folds for each of 8 candidates, totalling 24 fits
Wall time: 11min 53s
```

```
In [99]: gs
```

Out[99]:



In [100...

gs.best_params_

Out[100...

```
{'regressor': RandomForestRegressor(),
 'regressor__max_depth': 15,
 'regressor__n_estimators': 20}
```

In [101...

gs.score(X_test_pipe,y_test_pipe)

Out[101...

-45.8964777603188

In [102...

pred_y = gs.predict(X_test_pipe)

In [103...

r2_score(y_test_pipe,pred_y)

Out[103...

0.38487937564878716

So After using Gridsearch I found that RandomForestRegressor is the best model with max_depth 15 and n_estimators 20 which shows R2 score 0.3848 and negative mean square error is -45.89.

In []:

To enhance the RandomForestRegressor model's performance, consider tuning hyperparameters such as max_depth and n_estimators .

- **max_depth** : Determines the maximum depth of each tree in the forest. Higher values can lead to overfitting, so finding an optimal depth is crucial for balancing model complexity.
- **n_estimators** : Defines the number of trees in the forest. While increasing this parameter can enhance performance, excessively high values might not significantly improve results and can increase computational load.

So i should use more number max_depth and n_estimators to improve my model performance. But its taking too much time tahts why I have used only 15 as max_debth and

n_estimators as 20.

In []:

In [104...

```
"""This is the pipeline To drop specific columns """
drop_transformer = ColumnTransformer(transformers=[('drop_columns', 'drop', drop_featu
```

In [105...

```
# drop_transformer = ColumnTransformer(transformers=[('drop_columns', 'drop', drop_fea

# Creating pipelines for numerical and categorical features
numerical_pipeline = Pipeline(steps=[
#     ('outlier_removal', remove_outliers1(remove_outliers)),
    ('imputer', SimpleImputer(strategy='mean')) ,
    ('stdscaler', StandardScaler())
])

categorical_pipeline = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='most_frequent')), # Filling missing values
    ('onehot', OneHotEncoder(handle_unknown='ignore')) # One-hot encoding categoric
])

# Creating a ColumnTransformer to apply the appropriate pipeline to each type of fea
col_transformer = ColumnTransformer(transformers=[
    ('drop_columns', 'drop', drop_features_1),
    ('numerical', numerical_pipeline, Numerical_features_2),
    ('categorical', categorical_pipeline, categorical_features_1),
#     ('scale', StandardScaler())
], remainder='drop')

pipe_2pp = Pipeline(steps=[

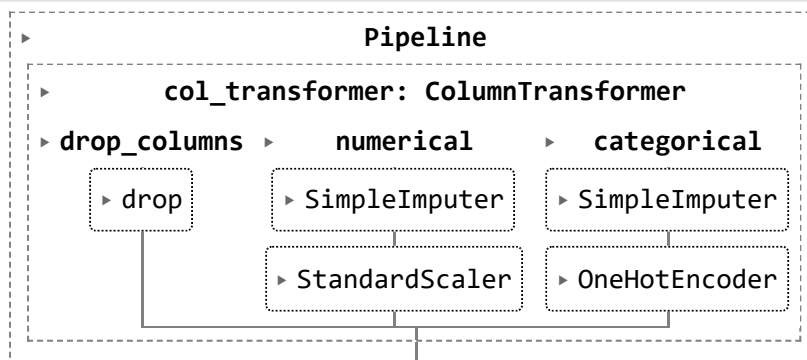
#     ('drop_transformer', drop_transformer),
#         ('col_transformer', col_transformer),

#     ('pca', PCA()),
#
#     ])
)
```

In [106...

pipe_2pp

Out[106...



In []:

In [107...

```
# Initialize the estimators
clf1 = RandomForestRegressor()
clf2 = Lasso()

clf3 = LinearRegression()
clf4 = Ridge()
clf5 = ElasticNet()

clf6 = DecisionTreeRegressor()
```

In [108...

```
# Initiaze the hyperparameters for each dictionary

#hyperparameters for RandomForestRegressor
param1 = {}
param1['regressor__n_estimators'] = [250]
param1['regressor__max_depth'] = [50]
param1['regressor'] = [clf1]

#hyperparameters for Lasso

param2 = {}
param2['regressor__alpha'] = [0.1, 1, 10]
param2['regressor'] = [clf2]

#hyperparameters for LinearRegression

param3 = {}
param3['regressor'] = [clf3]

#hyperparameters for Ridge

param4 = {}
param4['regressor__alpha'] = [0.1, 1]
param4['regressor'] = [clf4]

#hyperparameters for ElasticNet

param5 = {}
param5['regressor__alpha'] = [0.1, 1]
param5['regressor__l1_ratio'] = [0.2, 0.5, 0.7]
param5['regressor'] = [clf5]

#hyperparameters for DecisionTreeRegressor

param6 = {}
param6['regressor__max_depth'] = [3, 6]
param6['regressor'] = [clf6]
```

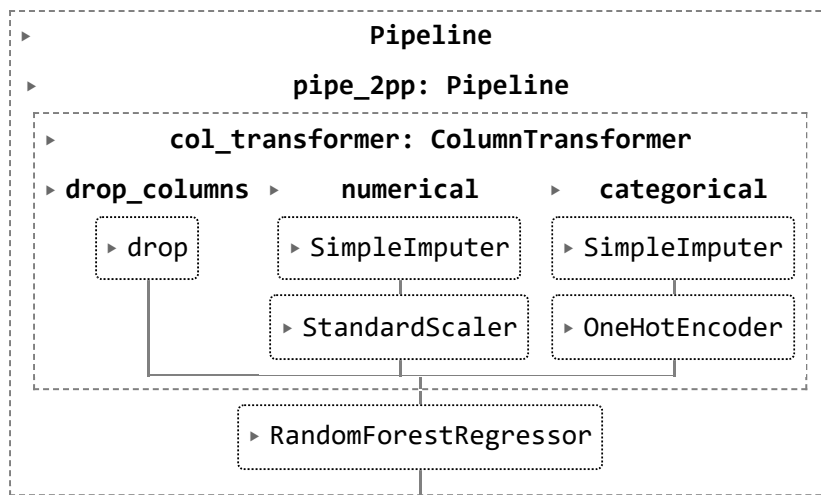
In [109...

```
pipeline = Pipeline(steps=[('pipe_2pp', pipe_2pp),
                           ('regressor', clf1)])
params = [param1, param2, param3, param4, param5, param6]
# params = [param1]
```

In [110...

```
pipeline
```

Out[110...



In [111...

```
# Gridsearchcv is used
# %%time
```

```
gs = GridSearchCV(pipeline, params, cv=3, n_jobs=-1, scoring='neg_root_mean_squared_
```

In [112...

```
%%time
```

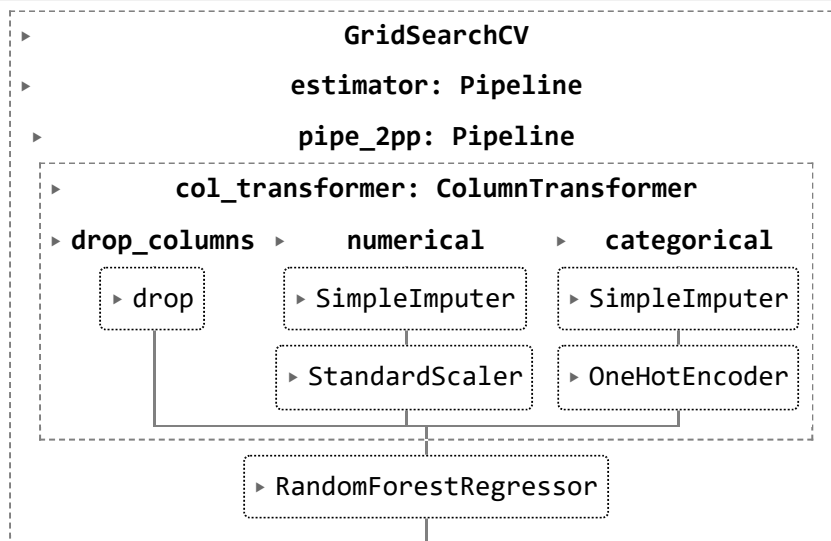
```
gs_fit = gs.fit(X_train_pipe, y_train_pipe)
```

Fitting 3 folds for each of 15 candidates, totalling 45 fits
Wall time: 16min 54s

In [113...

```
gs
```

Out[113...



In [114...

```
gs.best_params_
```

Out[114...

```
{'regressor': RandomForestRegressor(),
 'regressor__max_depth': 50,
 'regressor__n_estimators': 250}
```

In [115...

```
gs.score(X_test_pipe,y_test_pipe)
```

Out[115...

```
-39.12102044360718
```

In [116...

```
pred_y = gs.predict(X_test_pipe)
```

In [117...

```
r2_score(y_test_pipe, pred_y)
```

Out[117...

```
0.5530880963854016
```

So After using Gridsearch I found that RandomForestRegressor is the best model with max_depth 50 and n_estimators 250 which shows R2 score 0.5530 and negative mean square error is -39.12 [with out PCA]

As without PCA the RMSE reduced from 45 to 39 and the speed of the model training in improved. So without PCA my model works better.

In []:

In []: